

Anton Aspidov

AI / Full-stack Product Engineer — LLM systems that hold up in production

Saint Petersburg, Russia · Remote, UTC+3 · English B1, Russian native
anthonypriceone@gmail.com · Telegram @AnthonyPriceOne · linkedin.com/in/anton-aspidov
Portfolio with demos: portfolio-site.anthony-priceone.workers.dev

Profile

I design and ship LLM products end to end — data model, backend, queues, UI, integrations, deploy — and run them in production. I build with AI coding agents at team speed, inside a quality contour I designed (specs, 22 hard-capped gates, a doctor that audits the gates themselves); architecture, contracts and the final review stay mine. That is how two client services reached production in 8 and 14 days.

What I bring

- **LLMs under control, not just connected.** Behavioral contracts as versioned specs: hard rules block a message, soft rules force a rewrite, a live switch sets checks to full, warn-only or off without a restart. Autonomy is earned by measurement, not trust: shadow runs on real mail score every draft; auto-send stays off until replay clears the bar.
- **Speed without shortcuts.** Size-capped spec → plan → verify slices; gates that block the merge (strict typing, import-layer contracts, ratchets that only tighten); every live failure becomes a test. The same contour runs in 7 repos.
- **Cost engineering for paid APIs.** Pre-run estimates, fail-closed quota checks, caching and per-request cost ledgers: live probes showed an estimate built from the vendor's docs was 9× too low, and a three-stage filtering funnel cut spend by 37% on a measured run.
- **Evaluation discipline.** Held-out sets with a ledger of every look, predictions written before the run, retrieval calibrated against a public benchmark (BM25 nDCG@10 0.666 vs a 0.665 anchor), golden sets behind LLM judges.
- **Private, local AI.** Ollama-based agents that never leave the laptop: a browsing research agent and a voice interviewer whose first word lands ~1.3 s after the end of your turn is detected.

Selected work

- **LinkBuilder — outreach CRM** (*ran in production, sole engineer*). One operator ran what took a team of five: competitor mining, site-quality gates, mass outreach, AI-drafted replies under versioned contracts, payment requests, and monitoring of 16,000+ placed links. 11,000+ backend and 700 frontend tests, 137 migrations, 20+ workers.
- **SEO case-study platform** (*production, client under NDA*). Ahrefs history → classification → one-page PDF case studies. Live data in 4 days, production on day 14 against a 4–4.5-week estimate; 959 tests; paid runs commit after every task, so a crash loses only in-flight requests.
- **Prospecting & outreach platform** (*donor search live, client under NDA*). A GPT-5-mini site judge that quotes the page for every verdict, email-masked letter rewriting and reply-price extraction; eval gates on two of three. Live in 8 days; 100 PRs and 2,200+ tests by day 10.
- **G connect — MCP server** for Google Sheets and Drive, so Claude Code and Cursor agents can edit safely: 6 tools, 34 policy rules enforced in code, preview-first writes confirmed by plan hash, snapshot undo. Built in 8 days, 400+ tests.
- **Local Web Agent · Voice Interview Coach · RAG over a rulebook.** An on-device agent that browses real sites and backs its answers with quotes checked against the page (451 tests, open source); an offline voice interviewer with a per-sentence contract; retrieval with zero uncited answers.

Experience

Bookmaker Ratings — AI Automation Engineer / Internal Tools Builder · Aug 2021 — Present · Remote

- Built and owned LinkBuilder in production: architecture, refactor waves, human-in-the-loop design, and a handover document that describes the live state of the system, not its design.
- Built content automation: ~5× faster content preparation, ~120 hours a month saved. SEO background: grew organic traffic from 30–40K to ~180K monthly visits (Brazil, Portugal).

Skills

Backend: Python, FastAPI, asyncio, SQLAlchemy 2, Pydantic, Alembic, PostgreSQL, Redis/RQ, Node.js · **Frontend:** React, TypeScript, Vite, Mantine, TanStack Query · **AI:** OpenAI API, Ollama, tool calling, MCP, behavioral contracts, HITL, evals, RAG, ASR/TTS · **Delivery:** Claude Code, Cursor, pre-commit gates, GitHub Actions, Docker, Playwright, pytest, Vitest, mutation testing

Education: Saint Petersburg State University of Service and Economics, 2014